



Graph Algorithms

From Paths to Optimization



Evís Plaku

Why study graph algorithms?

I know a shortcut



...BFS taking every possible route

A* Algorithm




watching everyone else
explore useless paths

Shortest path algorithms

Introduction to path finding problems

Finding the shortest or optimal path in a graph is essential in many applications

- **Graph representation:** nodes and edges represent locations and connections in real-world scenarios
- **Shortest path:** goal is to find the quickest route from a starting point to a destination
- Used in **navigation systems, network routing, robotics, and AI pathfinding**




Unweighted graphs: BFS for shortest path




BFS explores nodes level by level, guaranteeing the shortest path in unweighted graphs

- Explores all neighbors at current depth before moving to the next level
- **No weights:** all edges are treated equally, making BFS ideal for unweighted graphs
- BFS ensures the **shortest path** by visiting nodes in increasing distance from the source




Single source shortest path problem



Find the shortest path from one starting node to all other nodes in a graph

- **Problem:** given a source node, find the shortest path to all other nodes
- Applies to both **directed** and **undirected graphs**, weighted or unweighted.
- Critical in **network routing, navigation systems**, and **AI game pathfinding**



Dijkstra's algorithm: intuition



Dijkstra's algorithm finds the shortest path by greedily selecting the closest node at each step


- **Greedy:** select the **nearest unvisited node** and update shortest distances iteratively
- **Priority queue:** manages nodes to ensure the closest node is processed first
- Dijkstra works when all edge weights are **non-negative** to ensure correctness

Edsger Dijkstra (1930 - 2002)



Perfecting oneself is as much
unlearning as it is learning

Dijkstra's algorithm: intuition



Dijkstra's algorithm explores nodes in order of closest distance, updating paths until all nodes are visited

Dijkstra is optimal for graphs with non-negative weights

Dijkstra's algorithm

```
1 function Dijkstra(Graph, source):  
2  
3   for each vertex v in Graph.Vertices:  
4     dist[v] ← INFINITY  
5     prev[v] ← UNDEFINED  
6     add v to Q  
7     dist[source] ← 0  
8  
9   while Q is not empty:  
10    u ← vertex in Q with minimum dist[u]  
11    remove u from Q  
12  
13    for each neighbor v of u still in Q:  
14      alt ← dist[u] + Graph.Edges(u,  
15      v)  
16      if alt < dist[v]:  
17        dist[v] ← alt  
18        prev[v] ← u  
19  return dist[], prev[]
```

- **Input:** a graph $G = (V, E)$ consisting of a set of nodes and a set of *weighted edges*
- **Set up a map of tentative distances**
- Every node is assumed unreachable at first (distance: ∞)
- Only the source has a known distance (0)
- As the algorithm runs, we will update $dist[v]$ with shorter paths and **track** $prev[v]$ to reconstruct the optimal route later

Dijkstra's algorithm

```
1 function Dijkstra(Graph, source):
2
3     for each vertex v in Graph.Vertices:
4         dist[v] ← INFINITY
5         prev[v] ← UNDEFINED
6         add v to Q
7     dist[source] ← 0
8
9     while Q is not empty:
10         u ← vertex in Q with minimum dist[u]
11         remove u from Q
12
13         for each neighbor v of u still in Q:
14             alt ← dist[u] + Graph.Edges(u,
15 v)
16             if alt < dist[v]:
17                 dist[v] ← alt
18                 prev[v] ← u
19     return dist[], prev[]
```

- **Select closest node:** pick the unvisited node with the smallest known distance
- **Mark as visited:** remove it from the queue; its shortest path is now finalized
- Continue until all nodes have been visited and processed

Dijkstra's algorithm

```
1 function Dijkstra(Graph, source):
2
3   for each vertex v in Graph.Vertices:
4     dist[v] ← INFINITY
5     prev[v] ← UNDEFINED
6     add v to Q
7   dist[source] ← 0
8
9   while Q is not empty:
10    u ← vertex in Q with minimum dist[u]
11    remove u from Q
12
13    for each neighbor v of u still in Q:
14      alt ← dist[u] + Graph.Edges(u,
15      v)
16      if alt < dist[v]:
17        dist[v] ← alt
18        prev[v] ← u
19
20  return dist[], prev[]
```


- **Explore neighbors:** for each unvisited neighbor v of the current node u , calculate a new possible distance
- **Update if shorter:** if the new distance alt is shorter than known distance $dist[v]$ update $dist[v]$
- **Track path:** set $prev[v] \mapsto u$ to remember that the shortest path to v goes through u
- **Return results:** output the shortest distance and paths from source to all other nodes

Dijkstra vs BFS: when to use each

BFS




- Best used for unweighted graphs
- Shortest path = fewest edges
- Simpler and faster



Dijkstra

- Weighted graphs with non-negative edge
- Shortest path = minimal total weight
- More flexible, handles real-world distances




Negative weights: why Dijkstra fails?



Dijkstra assumes once a node's shortest path is found, it won't be improved: this breaks with negative edges

- **Greedy assumption:** Dijkstra finalizes a node too early, assuming no shorter path will appear later
- **Negative edge impact:** a cheaper path might appear **after** visiting a node, but Dijkstra won't revisit it
- Incorrect results: this can lead to **wrong shortest paths** in graphs with negative edge weights




Negative weights examples:
profit in financial trading,
rewards in games, savings in
routing, or task penalties

Bellman – Ford algorithm

Bellman-Ford finds shortest paths, handles negative weights, and detects negative cycles

- Works even when edges have negative values



- Identifies cycles where total weight is negative
- **Slower than Dijkstra:** but necessary for graphs with negative edges

Bellman – Ford algorithm intuition

```
function BellmanFord(vertices, edges, source):
```

```
// Initialize distance and predecessor arrays
```

```
distance := [ $\infty$ ] * n
```

```
predecessor := [null] * n
```

```
distance[source] := 0
```

```
// Step 1: Relax edges ( $|V| - 1$ ) times
```

```
for i from 1 to  $|V| - 1$ :
```

```
    for each edge (u, v, w) in edges:
```

```
        if distance[u] + w < distance[v]:
```

```
            distance[v] := distance[u] + w
```

```
            predecessor[v] := u
```

```
// Step 2: Check for negative-weight cycles
```

```
for each edge (u, v, w) in edges:
```

```
    if distance[u] + w < distance[v]:
```

```
        // Negative cycle detected, find the cycle
```

```
        visited := [false] * n
```

```
        visited[v] := true
```

```
        while not visited[u]:
```

```
            visited[u] := true
```

```
            u := predecessor[u]
```

```
        cycle := [u]
```

```
        v := predecessor[u]
```

```
        while v != u:
```

```
            cycle := [v] + cycle
```

```
            v := predecessor[v]
```

```
        error "Negative-weight cycle detected", cycle
```

```
return distance, predecessor
```

- **Input:** list of vertices and edges in the graph and the starting vertex

Initialization

- **Distance:** set the initial distance to all vertices as infinity (∞) except the source (0)
- **Predecessor:** set all predecessors to **null** as no paths are established yet
- **Source:** the distance from the source to itself if (0)

Bellman – Ford algorithm intuition

```
function BellmanFord(vertices, edges, source):  
    // Initialize distance and predecessor arrays  
    distance := [ $\infty$ ] * n  
    predecessor := [null] * n  
    distance[source] := 0
```

```
// Step 1: Relax edges ( $|V| - 1$ ) times  
for i from 1 to  $|V| - 1$ :  
    for each edge (u, v, w) in edges:  
        if distance[u] + w < distance[v]:  
            distance[v] := distance[u] + w  
            predecessor[v] := u
```

```
// Step 2: Check for negative-weight cycles  
for each edge (u, v, w) in edges:  
    if distance[u] + w < distance[v]:  
        // Negative cycle detected, find the cycle  
        visited := [false] * n  
        visited[v] := true  
        while not visited[u]:  
            visited[u] := true  
            u := predecessor[u]  
        cycle := [u]  
        v := predecessor[u]  
        while v != u:  
            cycle := [v] + cycle  
            v := predecessor[v]  
        error "Negative-weight cycle detected", cycle
```

```
return distance, predecessor
```

Repeatedly update the shortest path estimates by "relaxing" all edges

- **Relax an edge:** for each edge (u, v, w) if the path from u to v is shorter than current known path, update both the distance and predecessor of v
- **Why $|V| - 1$ times?** In a graph with $|V|$ vertices, the longest path can have at most $|V| - 1$ edges
- Repeating this process ensures shortest paths are found

Bellman – Ford algorithm intuition

```
function BellmanFord(vertices, edges, source):  
    // Initialize distance and predecessor arrays  
    distance := [ $\infty$ ] * n  
    predecessor := [null] * n  
    distance[source] := 0
```

```
// Step 1: Relax edges ( $|V| - 1$ ) times  
for i from 1 to  $|V| - 1$ :  
    for each edge (u, v, w) in edges:  
        if distance[u] + w < distance[v]:  
            distance[v] := distance[u] + w  
            predecessor[v] := u
```

```
// Step 2: Check for negative-weight cycles  
for each edge (u, v, w) in edges:  
    if distance[u] + w < distance[v]:  
        // Negative cycle detected, find the cycle  
        visited := [false] * n  
        visited[v] := true  
        while not visited[u]:  
            visited[u] := true  
            u := predecessor[u]  
        cycle := [u]  
        v := predecessor[u]  
        while v != u:  
            cycle := [v] + cycle  
            v := predecessor[v]  
        error "Negative-weight cycle detected", cycle
```

```
return distance, predecessor
```

Repeatedly update the shortest path estimates by "relaxing" all edges

- **Relax an edge:** for each edge (u, v, w) if the path from u to v is shorter than current known path, update both the distance and predecessor of v
- **Why $|V| - 1$ times?** In a graph with $|V|$ vertices, the longest path can have at most $|V| - 1$ edges
- Repeating this process ensures shortest paths are found

Bellman – Ford algorithm intuition

```
function BellmanFord(vertices, edges, source):
    // Initialize distance and predecessor arrays
    distance := [∞] * n
    predecessor := [null] * n
    distance[source] := 0

    // Step 1: Relax edges (|V| - 1) times
    for i from 1 to |V| - 1:
        for each edge (u, v, w) in edges:
            if distance[u] + w < distance[v]:
                distance[v] := distance[u] + w
                predecessor[v] := u

    // Step 2: Check for negative-weight cycles
    for each edge (u, v, w) in edges:
        if distance[u] + w < distance[v]:
            // Negative cycle detected, find the cycle
            visited := [false] * n
            visited[v] := true
            while not visited[u]:
                visited[u] := true
                u := predecessor[u]
            cycle := [u]
            v := predecessor[u]
            while v != u:
                cycle := [v] + cycle
                v := predecessor[v]
            error "Negative-weight cycle detected", cycle

    return distance, predecessor
```

Detect if any edge can still be relaxed,
indicating a negative-weight cycle

- If any edge (u, v, w) can still be relaxed, it suggest a negative – weight cycle exists
- If such cycle exists, backtrack from v to find the vertices involved in the cycle
- The *visited* array is used to mark all visited nodes during backtrack to avoid infinite loops and detect where the cycle starts

Bellman – Ford algorithm intuition

```
function BellmanFord(vertices, edges, source):
    // Initialize distance and predecessor arrays
    distance := [∞] * n
    predecessor := [null] * n
    distance[source] := 0

    // Step 1: Relax edges (|V| - 1) times
    for i from 1 to |V| - 1:
        for each edge (u, v, w) in edges:
            if distance[u] + w < distance[v]:
                distance[v] := distance[u] + w
                predecessor[v] := u

    // Step 2: Check for negative-weight cycles
    for each edge (u, v, w) in edges:
        if distance[u] + w < distance[v]:
            // Negative cycle detected, find the cycle
            visited := [false] * n
            visited[v] := true
            while not visited[u]:
                visited[u] := true
                u := predecessor[u]
            cycle := [u]
            v := predecessor[u]
            while v != u:
                cycle := [v] + cycle
                v := predecessor[v]
            error "Negative-weight cycle detected", cycle

    return distance, predecessor
```

Track and build the cycle

- Once a cycle is detected, trace the cycle path using the *predecessor* links
- Keep adding nodes to *cycle* array until you loop back to the starting node *u*
- Once the cycle is fully traced, report it as a negative-weight cycle and throw an error
- The algorithm returns the **shortest path distances** and the **predecessors** for each vertex or reports **negative-weight cycle** if detected

Dijkstra vs Bellman – Ford

Feature	Dijkstra	Bellman – Ford
Edge weights	Works only with non-negative weights	Handles negative weights
Time Complexity	$O(V \log V)$ with a priority queue	$O(V * E)$ (slower)
Negative cycles	Does not detect negative cycles	Detects negative-weight cycles
Efficiency	More efficient for graphs with positive weights	Suitable for graphs with negative weights
Ideal Use Case	Shortest paths in positive weight graphs	Shortest paths with negative weights

Heuristic search

Informed search through heuristics



Heuristics guide search algorithms by estimating how close a state is to the goal


- **Uninformed search** explores blindly. It has no sense of direction or how far the goal is
- **Heuristics add intelligence** by using domain knowledge to rank possible paths
- Solutions are found faster and more efficiently



Informed search through heuristics

A heuristic is a function that estimates the cost from a node to the goal

- **Notation:** $h(n)$ is the estimated cost from node n to the goal
- Design: a good heuristic is **admissible** (never overestimates) and efficient to compute
- Example of admissible and efficient heuristic: **straight-line distance** (Euclidean distance) to the goal





A* search: intuition



A* finds the shortest path by combining actual cost so far and estimated cost to goal

- Select nodes using the best balance of **progress** and **promise**
- **Efficient and optimal** when the heuristic is admissible and consistent



A* was invented by researchers working on Shakey the Robot's path planning

A* search key components

Actual cost $g(n)$

- Represents the **cost** to reach node n from the **start** node
- **Cumulative** cost accumulated by following the path from start to n
- Helps ensure the algorithm chooses the shortest path


Heuristic estimate $h(n)$

- Represents the **estimated cost** from node n to the **goal**
- Uses domain knowledge (like straight-line distance) to guide search
- A good heuristic is admissible, meaning it never overestimates the true cost

$$f(n) = g(n) + h(n)$$

- Total estimated cost combines the actual cost with heuristic estimate
- Helps the algorithm balance between **exploring** already explored paths (low $g(n)$) and promising paths (low $h(n)$).
- Selects the node with the lowest $f(n)$

A* algorithm



Key: green: start; blue: goal; orange: visited

An example of an A* algorithm in action where nodes are cities connected with roads and $h(x)$ is the straight-line distance to the target point

A* algorithm

```
function A*(start, goal):
    open_list := priority queue with start node
    closed_list := empty set

    g[start] := 0                # Cost from start to start
    h[start] := heuristic(start) # Estimated cost from start to goal
    f[start] := g[start] + h[start] # Total estimated cost

    while open_list is not empty:
        current := node in open_list with the lowest f[current]
        if current == goal:
            return reconstruct_path(current) # Path from start to goal

        remove current from open_list
        add current to closed_list

        for each neighbor of current:
            if neighbor in closed_list:
                continue

            tentative_g := g[current] + distance(current, neighbor)

            if neighbor not in open_list:
                add neighbor to open_list
            else if tentative_g >= g[neighbor]:
                continue

            g[neighbor] := tentative_g
            h[neighbor] := heuristic(neighbor)
            f[neighbor] := g[neighbor] + h[neighbor]
            set neighbor's parent to current

    return failure # If no path found
```

- The **open list** stores nodes that are yet to be evaluated (implemented as priority queue)
- **Start node** is initially added to this list because it's the first node to explore
- The **closed list** tracks nodes that have already been fully evaluated (they won't be revisited)
- $g(n)$ represents the **actual cost** to reach node n from the start (0 for start)
- $h(n)$ is the **heuristic estimate** of the cost to reach the goal from node n (e.g., straight line)
- $f(n)$ is the total estimated cost, combining $g(n)$ and $h(n)$

A* algorithm

```
function A*(start, goal):
    open_list := priority queue with start node
    closed_list := empty set

    g[start] := 0                # Cost from start to start
    h[start] := heuristic(start) # Estimated cost from start to goal
    f[start] := g[start] + h[start] # Total estimated cost

    while open_list is not empty:
        current := node in open_list with the lowest f[current]
        if current == goal:
            return reconstruct_path(current) # Path from start to goal

        remove current from open_list
        add current to closed_list

        for each neighbor of current:
            if neighbor in closed_list:
                continue

            tentative_g := g[current] + distance(current, neighbor)

            if neighbor not in open_list:
                add neighbor to open_list
            else if tentative_g >= g[neighbor]:
                continue

            g[neighbor] := tentative_g
            h[neighbor] := heuristic(neighbor)
            f[neighbor] := g[neighbor] + h[neighbor]
            set neighbor's parent to current

    return failure # If no path found
```

- The loop continues as long as there are nodes left to explore
- Select the node current from the **open list** that has the **lowest $f(n)$** value (explore most promising node first)
- If the **current node** is the goal, **return the path** from the start to the goal
- After exploring the **current node**, remove it from the **open list** since it's now fully evaluated
- Add the **current node** to the **closed list** to indicate that it's been fully evaluated

A* algorithm

```
function A*(start, goal):
    open_list := priority queue with start node
    closed_list := empty set

    g[start] := 0                # Cost from start to start
    h[start] := heuristic(start)  # Estimated cost from start to goal
    f[start] := g[start] + h[start] # Total estimated cost

    while open_list is not empty:
        current := node in open_list with the lowest f[current]
        if current == goal:
            return reconstruct_path(current) # Path from start to goal

        remove current from open_list
        add current to closed_list

        for each neighbor of current:
            if neighbor in closed_list:
                continue

            tentative_g := g[current] + distance(current, neighbor)

            if neighbor not in open_list:
                add neighbor to open_list
            else if tentative_g >= g[neighbor]:
                continue

            g[neighbor] := tentative_g
            h[neighbor] := heuristic(neighbor)
            f[neighbor] := g[neighbor] + h[neighbor]
            set neighbor's parent to current

    return failure # If no path found
```

- For each **neighbor** of the **current node**, the algorithm considers the possibility of reaching that neighbor from the current node
- If the **neighbor** has already been fully evaluated (i.e., it is in the **closed list**), skip it and move to the next neighbor
- *tentative_g* calculates the new cost to reach the **neighbor** from the start node by adding the cost of reaching the current node and the cost to move from current node to neighbor
- If the neighbor is undiscovered, add it to the open list. Otherwise, skip if cost is higher

A* algorithm

```
function A*(start, goal):
    open_list := priority queue with start node
    closed_list := empty set

    g[start] := 0                # Cost from start to start
    h[start] := heuristic(start) # Estimated cost from start to goal
    f[start] := g[start] + h[start] # Total estimated cost

    while open_list is not empty:
        current := node in open_list with the lowest f[current]
        if current == goal:
            return reconstruct_path(current) # Path from start to goal

        remove current from open_list
        add current to closed_list

        for each neighbor of current:
            if neighbor in closed_list:
                continue

            tentative_g := g[current] + distance(current, neighbor)

            if neighbor not in open_list:
                add neighbor to open_list
            else if tentative_g <= g[neighbor]:
                continue

            g[neighbor] := tentative_g
            h[neighbor] := heuristic(neighbor)
            f[neighbor] := g[neighbor] + h[neighbor]
            set neighbor's parent to current

    return failure # If no path found
```

- Update the **g value** of the **neighbor** with the newly calculated cost from the start
- Update the **h value** of the **neighbor** with its heuristic estimate to the goal
- Update the **f value** of the **neighbor** as the sum of its **g** and **h** values
- Set the **current node** as the **parent** of the **neighbor** to trace the path later

The algorithm terminates
when the goal node is reached
or the open list is empty

A* applications


Pathfinding in Video Games




Navigation Systems



Robotics



AI in Puzzles



Autonomous Vehicles



Network Routing



There's no shortcut to success ...

except in graphs, where finding the shortest path is the goal!

